

# A model-based infrastructure for the specification and runtime execution of self-adaptive IoT architectures

Iván Alfonso<sup>1,2\*</sup>, Kelly Garcés<sup>1</sup>, Harold Castro<sup>1</sup> and Jordi Cabot<sup>2,3</sup>

<sup>1\*</sup>Department of Systems and Computing Engineering,  
Universidad de los Andes, Bogotá, Colombia.

<sup>2</sup>Universitat Obertat de Catalunya, Barcelona, Spain.

<sup>3</sup>ICREA, Barcelona, Spain.

\*Corresponding author(s). E-mail(s): [id.alfonso@uniandes.edu.co](mailto:id.alfonso@uniandes.edu.co);  
Contributing authors: [kj.garces971@uniandes.edu.co](mailto:kj.garces971@uniandes.edu.co);  
[hcastro@uniandes.edu.co](mailto:hcastro@uniandes.edu.co); [jordi.cabot@icrea.cat](mailto:jordi.cabot@icrea.cat);

## Abstract

To meet increasingly restrictive requirements and improve quality of service (QoS), Internet of Things (IoT) systems have embraced multi-layered architectures leveraging edge and fog computing. However, the dynamic and changing IoT environment can impact QoS due to unexpected events. Therefore, proactive evolution and adaptation of the IoT system becomes a necessity and concern. In this paper, we present a model-based approach for the specification and execution of self-adaptive multi-layered IoT systems. Our proposal comprises the design of a domain-specific language (DSL) for the specification of self-adaptive IoT architectures, and a framework to support and adapt the system at runtime. The code for the deployment of the IoT system and the execution of the runtime framework is automatically produced by our prototype code generator. Moreover, we also show and validate the extensibility of such DSL by applying it to the domain of underground mining. The complete infrastructure (modeling tool, generator and runtime components) is available in an online open source repository.

**Keywords:** Domain-specific language, Internet of Things, Self-adaptive system, Edge computing, Fog computing

# 1 Introduction

The ideas behind IoT have been especially embraced by industry in the so-called Industrial IoT (IIoT) or Industry 4.0. Currently billions of devices become connected with potential capabilities to sense, communicate, and share information about their environment. Traditional IoT systems rely on cloud-based architectures, which allocate all processing and storage capabilities to cloud servers. Although cloud-based IoT architectures have advantages such as reduced maintenance costs and application development efforts, they also have limitations in bandwidth and communication delays [1]. Given these limitations, edge and fog computing have emerged with the goal of distributing processing and storage close to data sources (i.e. things). Today, developers tend to leverage the advantages of edge, fog, and cloud computing to design multi-layered architectures for IoT systems.

Nevertheless, creating such complex designs is a challenging task. Even more challenging is managing, and adapting IoT systems at runtime to ensure the optimal performance of the system while facing changes in the environmental conditions. IoT systems are commonly exposed to changing environments that induce unexpected events at runtime (such as unstable signal strength, latency growth, and software failures) that can impact its Quality of Service (QoS). To deal with such events, a number of runtime adaptations should be automatically applied. For example, *architectural adaptations* (such as auto-scaling and offloading tasks) to ensure system availability and QoS.

In this sense, a better support to define and execute complex IoT systems and their (self)adaptation rules to semi-automate the deployment and evolution process is necessary [2]. A usual strategy when it comes to modeling complex domains is to develop a domain-specific language (DSL) for that domain [3]. In short, a DSL offers a set of abstractions and vocabulary closer to the one already employed by domain experts. Nevertheless, current approaches for modeling IoT do not typically cover multi-layered architectures [4–7] and even less include a sublanguage to ease the definition of the dynamic rules governing the IoT system.

Our research aims at overcoming this situation by presenting a model-based infrastructure for the specification and runtime execution of multi-layered self-adaptive IoT architectures. Our proposal combines a DSL for the specification of static and dynamic aspects of this type of systems together with a runtime infrastructure and a code-generator able to semi-automate their deployment and runtime monitoring and adaptation.

This work is an extension of our study presented in [8], in which we proposed a first version of a DSL for IoT systems and a proof of concept of a code generator for the deploying of the initial configuration of the modeled IoT system using K3S<sup>1</sup> (Kubernetes distribution built for IoT and edge computing). The current work extends this previous contribution to the following aspects:

---

<sup>1</sup><https://k3s.io/>



- **Metamodel improvement:** we have enhanced the metamodel to support modeling new DSL concepts such as sensor threshold values, publish/subscribe messaging and data persistence for containers.
- **Runtime support:** we have developed a framework based on the MAPE-K [9] loop to automatically monitor and self-adapt the IoT system by executing the adaptation rules modeled with the DSL.
- **Code generator enhancements:** in addition to generating the YAML manifests for the deployment of the IoT system container-based applications, we can now also generate the code required to support the execution of the system at runtime (including the code for infrastructure monitoring and system management tools).
- **DSL extension for the mining industry:** we propose an extension of our DSL focused on the modeling and operation of IoT systems in the underground mining industry, highlighting the ability to reuse the DSL in other domains.
- **Empirical evaluations:** we have designed and conducted two empirical experiments to validate the expressiveness and usability of our DSL extended to the mining domain. One experiment focuses on device layer adaptations and the other one on architectural adaptations.

The remainder of the paper is organized as follows: Section 2 presents a running example to illustrate our approach. Section 3 introduces the DSL to specify self-adaptive IoT systems. Section 4 presents an extension of our DSL for IoT systems deployed in underground mining. In Section 5, we present the DSL implementation, the framework to support IoT systems at runtime, and code generation. In Section 6, we validate the usability of our DSL. Finally, the related work to this research is summarized in Section 7, and the conclusions and future work are presented in Section 8.


## 2 Running Example

We will use a Smart Building scenario as a main running example to better illustrate our approach.

Smart buildings seek to optimize different aspects such as ventilation, heating, lighting, energy efficiency, etc [10]. Adopting the concept of smart building, let's assume that a hotel company (*Hotel Beach*) wants to reduce fire risks by automating disaster management in its hotels. A fire alarm and monitoring system are implemented in each of the company's hotel. We will assume that all buildings (hotels) have three floors with two rooms each. Fig. 1 presents an overview of the 1st floor of a building. According to this, the infrastructure (device, edge, and cloud layers) of the company hotel IoT system are as follows:

- **Device layer.** Each room has a temperature sensor, a carbon monoxide (CO) gas sensor, and a fire water valve. Furthermore, an alarm is deployed on the lobby. Each sensor has a threshold measurement to activate the corresponding alarm, e.g., a person should not be continuously exposed to CO gas level of 50 parts per million (ppm) for more than 8 hours.

- **Edge layer.** In each room, an edge node receives the information collected by the sensors of the device layer and run a software container (C1 and C2) for analyzing sensor data in real time to check for the presence of smoke and generate an alarm state that activates the actuators. A fog node (linked to the edge nodes) is located in the 1st floor of the building. This node runs the C3 container (running App2, a machine learning model to predict fires on any of the building's floors), and C4 (running App3, in charge of receiving and distributing data, typically a Message Queuing Telemetry Transport (MQTT<sup>2</sup>) broker as we will see later on).
- **Cloud layer.** The cloud layer has a server or cloud node that runs the C5 container, a web application (App4) to display historical information of sensor data and of fire incidents in any of the hotel's property of the company.



**Fig. 1:** Overview of the smart building IoT system, first floor

In any IoT system, there are critical applications that should be available all the time. For example, the availability of the containers running App1 (real-time smoke monitoring to detect and alarm the presence of smoke) should be guaranteed. However, some environmental factors can generate unexpected events impacting system operation and services availability. In these cases, the IoT system must self-adapt to guarantee its operation. For instance, a fire could cause failures in the *edge-a1* node, then will be necessary to migrate the *C1* container to another suitable node.

<sup>2</sup>MQTT is a lightweight publish/subscribe messaging protocol

Our research addresses this type of system architectural adaptations, motivated by the detection of system irregularities (such as unavailability of applications or excessive CPU consumption of a node) or even environmental incidents that may induce failures (such as a fire). In addition to these architectural adaptations, we also address the definition of rules to meet the functional requirements of the system. For example, when one of the room sensors detects CO gas greater than 400 ppm, an of the alarms should be triggered. The next section shows how we can model all these concepts.

### 3 A DSL for the specification of multi-layered IoT systems

A DSL is defined by three core ingredients [3]: the abstract syntax (i.e., the concepts of the DSL and their relationships), the concrete syntax (the notation to represent those concepts), and its semantics which are hardly ever formalized but based on the shared understanding of the domain. In this chapter, we present our DSL for modeling the static (section 3.1) and dynamic (3.2) aspects of the IoT system.

#### 3.1 Modeling of the IoT Architecture

Multi-layered architectures have emerged to increase the flexibility of pure cloud deployments and help meet non-functional IoT system requirements [11]. In particular, edge computing and fog computing are solutions that propose to bring computation, storage, communication, control, and decision-making closer to IoT devices [12]. While fog computing take place on LAN-connected nodes usually close to end user devices, edge computing takes place on personal devices directly connected to sensors and actuators, or on gateways physically close to these [12, 13]. Sensors and actuators compose the device layer. Sensors generate information by monitoring physical variables (such as temperature, motion, and oxygen) and actuators (such as valves, fans, and alarms) are devices capable of transforming energy into activation of a process.

Edge and fog computing can leverage containerization as a virtualization technology [14]. Containers, compared to virtual machines, are lightweight, simple to deploy, support multiple architectures, have a short start-up time, and are suitable for dealing with the heterogeneity of edge and fog nodes. In terms of communications, asynchronous message-based architectures are typically adopted for cyber-physical and IoT systems that require high scalability [15]. The publisher/subscriber pattern and the MQTT protocol are becoming the standard for Machine-to-Machine (M2M<sup>3</sup>) communications [16], where messages are sent (by publishers) to a message broker server and routed to destination clients (subscribers).

The DSL enables the specification of all these concepts as part of a multi-layered IoT architectures.

---

<sup>3</sup>M2M is the communication and exchange of data between two machines or devices without human intervention.



Remark  
7



Remark  
5



Remark  
6

### 3.1.1 Abstract syntax

The abstract syntax of the DSL is commonly defined through a metamodel that represents the domain concepts and their relationships. Fig. 2 shows the metamodel that abstracts the concepts to define multi-layered IoT architectures. The sensors and actuators of the device layer are modeled using the *Sensor* and *Actuator* concepts that inherit from the *IoTDevice* concept. This generalization is restricted to be complete and disjoint. All *IoTDevices* have a connectivity type (such as Ethernet, Wi-Fi, ZigBee, or another). We also cover the concepts for specifying MQTT communication between IoT devices and nodes. Therefore, *IoTDevices* are publishers or subscribers to a topic MQTT broker specified by the relationship to the *Topic* concept. The gateway of an *IoTDevice* can be modeled through the *gateway* relationship with the *EdgeNode* concept. Via this gateway, the sensor can communicate with other nodes, e.g., the MQTT broker node. Additionally, the threshold value and unit of the monitored variable by a sensor can be represented through the attributes *threshold* and *unit*.

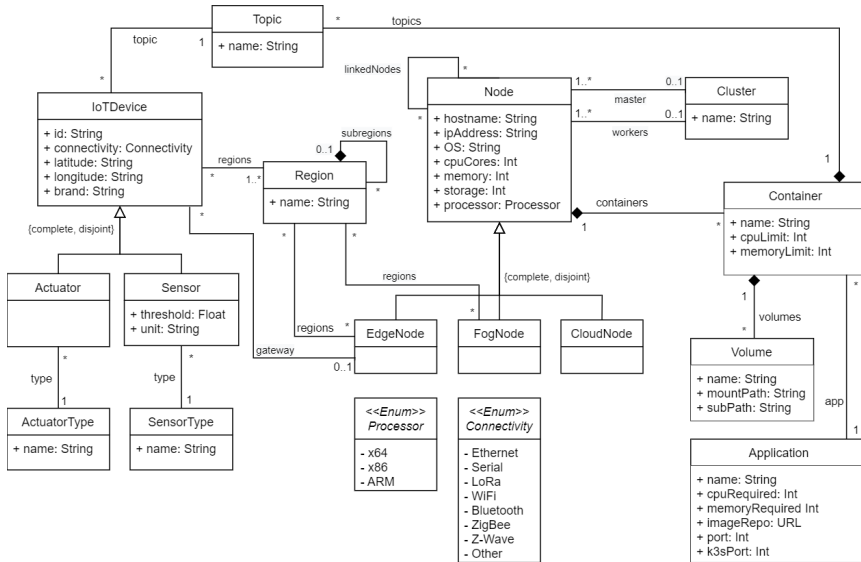
The location of *IoTDevices* can be specified through geographic coordinates (*latitude* and *longitude* attributes). Both *Sensors* and *Actuators* have a type represented by the concepts *SensorType* and *ActuatorType*. For instance, following the running example (Fig. 1), there are temperature and smoke type sensors, and there are valve and alarm type actuators.

Physical (or even virtual) spaces such as rooms, stairs, buildings, or tunnels can be represented by the concept *Region*. A *Region* can contain subregions (relationship *subregions* in the metamodel). For example, region *Floor1* (Fig. 1) contains subregions *Room1*, *Room2*, *Lobby*, and *Stairs*. *IoTDevices*, *EdgeNodes*, and *FogNodes* are deployed and are located in regions or subregions (represented by *regions* relationships in the metamodel). Back to the running example, the *edge-a1* node is located in the *RoomA1* region of *Floor1* of the *Hotel Beach*, while the *fog-f1* node is located in the *Lobby* region of *Floor1*.

Edge, fog and cloud nodes are all instances of *Node*, one of the key concepts of the metamodel. A node has the ability to host the software containers. Communication between nodes can be specified via the *linkedNodes* relationship, as we may want to indicate what nodes on a certain layer could act as reference nodes in another layer (e.g., what cloud node should be the first option for a fog node). Nodes can also be grouped in clusters that work together. The details of each node are expressed via attributes such as ip address (*ipAddress*), operating system (OS), number of cores in the processor (*cpuCores*), RAM memory (*memory*), storage capacity (*storage*), and processor type (*processor enum*).

A *Node* can host several software containers according to its capabilities and resources (primarily *cpuCores*, *memory*, and *storage*). The CPU and memory usage of a container can be restricted through *cpuLimit* and *memoryLimit* attributes. Each software container runs an application (represented by the concept *Application*) that has a minimum of required resources specified by the attributes *cpuRequired* and *memoryRequired*. The repository of the application image is specified through the *imageRepo* attribute, and the used

ports through *port* and *k3sPort* attributes. The container volumes and their paths (a mechanism for persisting data used and generated by containers) are represented by the *Volume* concept. Finally, the MQTT broker that receives and distributes the messages can also be specified and deployed in a software container, and the broker topics are represented by the *topics* relationship.



**Fig. 2:** Multi-layer IoT architecture metamodel

### 3.1.2 Concrete syntax

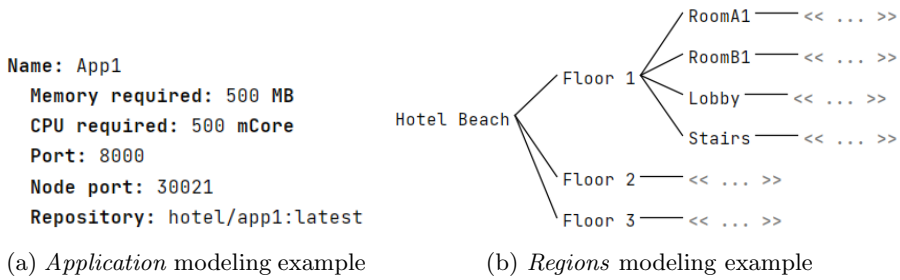
The concrete syntax refers to the type of notation (such as textual, or graphical) to represent the concepts of the metamodel. Graphical DSLs involve the development of models using graphic items such as blocks, arrows, axes, and so on. Textual DSLs involve modeling using configuration files and text notes. Though most DSLs employ a single type of notation, they could benefit from offering several alternative notations, each one targeting a different type of DSL user profile. This is the approach we follow here, leveraging the benefits of using a projectional editor.

Projectional editors enable the user’s editing actions to directly change the Abstract Syntax Tree (AST) without using a parser [17]. That is, while the editing experience simulates that of classical parsing-based editors, there is a single representation of the model stored as an AST and rendered in a variety of perspectives thanks to the corresponding projectional editors that can deal with mixed-language code and support various notations such as tables, mathematical formulas, or diagrams.

Indeed, we take advantage of MPS<sup>4</sup> projectional editors to define a set of complementary notations for the metamodel concepts. We blend textual, tabular, and tree view, depending on the element to be modeled. We next employ these notations to model our running example (2). More technical details on the implementation of our concrete notations are presented in Section 5.

### 3.1.3 Example scenario

First, Fig. 3a depicts the modeling of the IoT system application *App1*, including its technical requirements and repository address. Then, Fig. 3b shows the specification of the *Hotel Beach* regions, in particular those on *Floor 1* (four subregions: two *Rooms*, the *Lobby* and the *Stairs*).



**Fig. 3:** *Application* and *Regions* modeling example

Additionally, *IoT devices* can be modeled using a tabular notation. Fig. 4 shows the list of sensors and actuators located in the *RoomA1* region (we have removed the *latitude* and *longitude* device parameters from Fig. 4 to keep it readable). To specify the nodes of the system architecture, we propose a tabular notation as well (Fig. 5 shows *edge-a1* and *fog-f1* node modeling). The node description includes the layer it belongs to (edge, fog, or cloud), the hardware properties (such as memory and storage resources), the regions where it is located, and the application containers it hosts. Each container includes the specification of the application to be executed, CPU and memory limits, and volumes (e.g., the volume of the *C4* container for the configuration of the MQTT broker)

	Device	ID	Type	Unit	Threshold	Regions	Brand	Communic.	Gateway	Topic
1	Sensor	gas-a1	CO	ppm	50	RoomA1	Winsen	ZigBee	edge-a1	f1/rA1/smoke
2	Sensor	temp-a1	Temperature	°C	23	RoomA1	MLX	Z_Wave	edge-a1	f1/rA1/temp
3	Actuator	valve-a1	Valve	---	---	RoomA1	Bray	Serial	edge-a1	f1/rA1/valve

**Fig. 4:** IoT devices modeling example (tabular notation)

<sup>4</sup>Meta Programming System is an open-source language workbench developed by JetBrains



	Hostname	Layer	Properties	Regions	Linked nodes	Containers
1	edge-a1	Edge	Memory: 2000 MB Storage: 16000 MB CPU cores: 1 Core IP address: 192.168.10.1 Operating system: Raspbian Processor: ARM	RoomA1	fog-f1	* Name: C1 Application: App1 Memory limit: no limit CPU limit: no limit Volumes: << ... >>
2	fog-f1	Fog	Memory: 4000 MB Storage: 20000 MB CPU cores: 2 Cores IP address: 192.168.10.3 Operating system: Raspbian Processor: ARM	Lobby	cloud-hotel	* Name: C3 Application: App2 Memory limit: no limit CPU limit: no limit Volumes: << ... >> * Name: C4 Application: App3 Memory limit: no limit CPU limit: no limit Volumes: -> Name: mosquitto-config Mount path: /config/mqtt.conf Sub path: mosquitto.conf

Fig. 5: Nodes and containers modeling example (tabular notation)

## 3.2 Modeling Adaptation Rules

The dynamic environment of an IoT system generates unexpected changes that could affect the QoS. The system should be prepared to cope with unexpected changes and self-adapt to different situations. Because of this, our DSL enables the modeling adaptation rules of two types: rules to specify the self-adaptation of the system architecture, and rules that involve the operation of actuators to support system functionalities.

### 3.2.1 Abstract syntax

The metamodel representing the abstract syntax for defining the adaptation rules is presented in Fig. 6.

Every rule is an instance of *Rule* that has a triggering condition, which is an expression. We reuse an existing *Expression* sublanguage to avoid redefining in our language all the primitive data types and all the basic arithmetic and logic operations to manipulate them. Such Expression language could be, for instance, the Object Constraint Language (OCL), but to facilitate the implementation of the DSL later on, we directly reused the MPS *Baselanguage*. The metamodel extends the generic Expression concept by adding sensor and QoS conditions that can be combined also with all other types of expressions (e.g., numerical ones) in a complex conditional expression.

A *SensorEvent* represents the occurrence of an event resulting from the analysis of sensor data (e.g., the detection of dioxide carbon gas by the *gas-a1* sensor). Note that *SensorEvent* conditions can be linked to specific sensors or to sensor types to express conditions involving a group of sensors.

Similarly, the *QoSEvent* condition is a relational expression that represents a threshold of resource consumption or QoS metrics. This condition allows checking a *Metric* (such as Latency, CPU consumption, and others) on a specific node or a group of nodes belonging to a *Region* or *Cluster*. For example,

the condition  $\text{cpu}(\text{HotelBeach} - > \text{edge\_nodes}) > 50\%$  is triggered when the CPU consumption on the edge nodes of the *Hotel* exceeds 50%.

Moreover, we can define that the condition should be true over a certain period (to avoid firing the rule in reaction to minor disturbances) before executing the rule. Once fired, all or some of the actions are executed in order, depending on the *allActions* attribute. If set to false, only the number of *Actions* specified by the attribute *actionsQuantity* must be executed, starting with the first one in order and continuing until the required number of actions have been successfully applied.

We have classified the rules into two groups: *Architectural Adaptation Rules* composed of a *QoSEvent* condition and actions such as *Offloading*, *Redeploy*, or *Scaling*; and *Functional Rules* composed of a *SensorEvent* condition and *OperateActuator* actions. Note that they could be combined (e.g., a sensor event could trigger a scaling of apps to collect more data about the event or to double-check whether the sensor event is isolated or can be found in other locations).

The *Offloading* action consists in migrating a container from a source node to a destination node. This migration can be between nodes of different layers. The *container* relationship represents the container that will be offloaded. The target node is specified by the *targetNode* relationship. However, if the target node does not have the resources to host the container, a cluster or a group of nodes in a *Region* can be specified (*targetRegion* and *targetCluster* relationships) to offload the container. The *Scaling* action involves deploying replicas of an application. This application is represented by the *app* relationship, and the number of replicas to be deployed is defined by the *instances* attribute. The replicas of the application are deployed in one or several nodes of the system represented by the *targetNodes*, *targetCluster*, and *targetRegion* relationships. The *Redeployment* action consists in stopping and redeploying a container running on a node. The container to redeploy is indicated by the *container* relationship. Finally, the *OperateActuator* action is to control the actuators of the system (e.g., to activate or deactivate an alarm). The message attribute represents the message that will be published in the broker and interpreted by the actuator. This action can control actuators that require only one control value, such as On/Off.

### 3.2.2 Concrete syntax

These rules are specified thanks to a textual notation using as keywords the names of the metaclasses of the abstract syntax. The conditions follow the grammar of a relational expression with the use of mathematical symbols (such as  $<$ ,  $>$ , and  $=$ ) and logical operators (such as  $\&$  and  $\parallel$ ). The rule editor (see Section 5) offers a powerful autocomplete feature to guide the designer through the rule creation process.



Remark  
7



Remark  
7

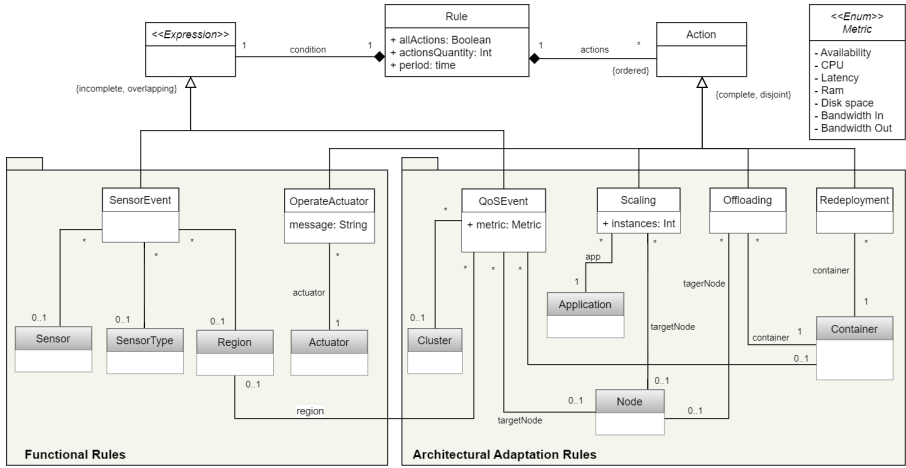


Fig. 6: IoT system adaptation rules metamodel

### 3.2.3 Example Scenario

We show how to use the rule's concrete syntax to model two example rules from the smart building example.

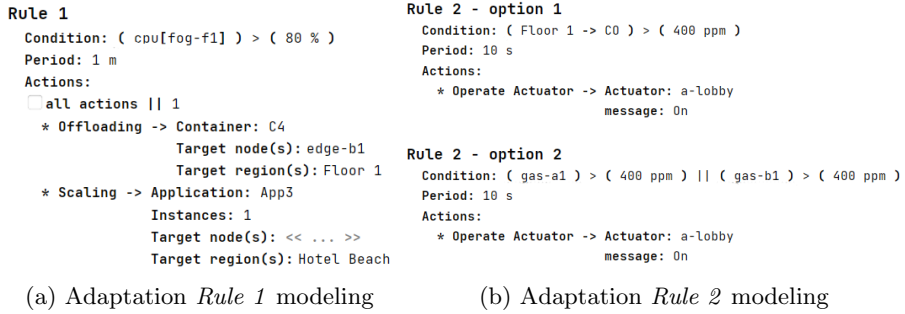
First, to guarantee the execution of the *C4* container deployed on the *fog-f1* node, we modeled the rule as shown in (Fig. 7a). This rule offloads the container *C4* hosted on node *fog-f1* to a nearby node (e.g., node *edge-b1*) when the CPU consumption exceeds 80% for one minute. If the *edge-b1* node does not have the necessary resources to host that new container (when the rule is activated), a Region (e.g., *Floor1*) can be specified so that a suitable node will be searched there. However, if this offloading action cannot be executed, for example, because in *Floor1* there is no node capable of hosting the container, then we must define a backup action. Therefore, we have modeled a second action (*Scaling*) to deploy a new container instance of the *App3* application on any of the nodes of the *Hotel Beach*. Although a list of actions can be specified, all or only a certain number of them may be performed. All actions in the list will be performed if the checkbox *all actions* is checked. Otherwise, only a defined number (specified after the checkbox) of actions will be performed. For this example adaptation rule (*Rule 1*), only one action (the first one, or the second one if the first one fails) will be executed.

Secondly, we model another rule (see Fig. 7b) to activate the alarm (a-lobby) when any gas sensors in the *Floor1* region (gas-a1 or gas-b1) detects a gas concentration greater than 400ppm for 10 seconds. The "On" message is published in the broker topic consumed by the actuator (*a-lobby* alarm). Note that there are two ways to model this rule. While Option 1 involves all CO type sensors on *Floor 1*, Option 2 directly involves both gas sensors.



Remark

21

**Fig. 7:** Adaptation rules

## 4 DSL Extension: Coal Underground Mining

Our DSL can be used as is to model any type of multi-layered IoT system. However, it has also been designed to be easily extensible so that we can further tailor it to specific types of IoT systems. As an example, we present an extension of our DSL to model IoT systems for underground mining as this is a key economic sector in the local region of one of the authors and there is a need for a better way to model these systems, e.g., for analysis of regulatory compliance.

The dynamic and hostile environment of the underground mining industry threatens the operation of IoT systems (e.g. by causing physical damage to the devices) implemented primarily to monitor and ensure the safety of workers. Explosive and toxic gases, risk of geotechnical failure, fire, high temperatures and humidity are some of the risks that threaten the proper functioning of the IoT monitoring system and, consequently, the safety of workers. Therefore, the system should cope with these unexpected changes by self-adapting to guarantee run-time operation (e.g. ensuring the availability of critical applications despite node failures).

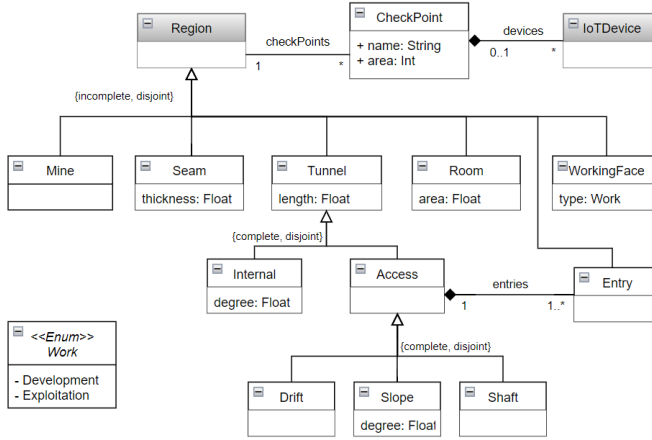
In addition to system self-adaptations, our approach could support functional requirements. For example, in Colombia, the safety regulation for underground mining works (decree 1886 of 2015) determines limits for the concentration of explosive and toxic gases. If any of these limits is exceeded, a series of actions/adaptations must be performed such as turning on alarms, activating the ventilation system, closing or opening ventilation ducts, and other tasks or rules that can be specified using our DSL.

To better cope with these scenarios, our extended DSL offers new modeling primitives (see Figure 8). All concepts that inherit from *Region* represent physical spaces within an underground coal mine such as tunnels, working faces, Entries, and rooms. A *Tunnel* can be *Internal* or *Access*. Each mine access tunnel (*Drift*, *Slope*, or *Shaft*) must have one or more entrances (represented by the *entries* relationship). Finally, checkpoints (areas of the mine where gases, temperature, oxygen, and airflow are monitored) are specified




Remark  
7

through the *CheckPoint* concept. Each *CheckPoint* could contain multiple *IoT-Devices* (sensors or actuators) represented by the *devices* relationship in the metamodel.



**Fig. 8:** Excerpt of the DSL extension metamodel

We offer a tree-based notation for modeling the relevant regions<sup>5</sup> that make up the mine structure. Figure 9 presents an example of the modeling of an underground mine containing two entries (*Entry A* and *Entry B*) in each of its inclined access tunnels (*Slope A* and *Slope B*), an internal tunnel (*Internal*), and a (*Room*) with two exploitation work fronts (*W-front 1* and *W-front 2*). The control points are modeled using textual notation, while the rest of the concepts to represent the IoT system and the adaptation rules are modeled following the concrete syntax presented in section 3.



**Fig. 9:** Mine structure modeling

At each mine control point, the airflow should be controlled by the fans. While very fast air currents can produce fires, very low air currents may not

<sup>5</sup>Note that our DSL is focused on the structure and rules governing the “behaviour” of the IoT system of the mine, it does not pretend to replace other types of 3D mine mining models

be efficient in dissipating gas concentrations. To involve control points directly in the adaptation rules, we have also modified our language. This extension of the DSL enables the modeling of conditions such as (*ControlPointA*  $\rightarrow$  *airFlow*) > (2m/s).

Other concepts such as the IoT system architecture and adaptation rules can be specified using the textual and tabular editors presented in Section 3.

## 5 Tool Support

In this Section, we describe the implementation and tool support we provide for both the implementation of our DSL, the core runtime framework in charge of monitoring and self-adapting the IoT systems at runtime, and the code-generator features that help you move from the system specification with our DSL to its deployment and runtime execution in our architecture.

### 5.1 DSL Tool Support

Our DSL is implemented using MPS, an open-source language workbench developed by JetBrains.

By building the DSL on top of MPS, we automatically get a projectional editor for the DSL with facilities to implement the different notations highlighted in the previous sections. The DSL editor is freely available in our repository<sup>6</sup>

#### 5.1.1 Language editor

To develop the modeling environment for the DSL, we had to define in MPS three core elements: structure, editors and constraints. The structure defines the abstract syntax of the DSL by defining all metamodel concepts. Projection editors define the AST code editing rules, and constraints define well-formedness rules in the editors. We have defined textual, tabular, and tree view editors by implementing the mbeddr<sup>7</sup> extension of MPS.

For example, Fig. 10 shows the definition of the tabular editor for modeling the *Sensor* concept. We have used the *partial table* command to define the table structure (cells, content, and column headers). By defining this editor, the user is enabled to model *Sensors* using a tabular notation as shown in Fig. 4.

#### 5.1.2 Constraints

Constraints restrict the relationships between nodes as well as the allowed values for properties. We have used this constraint mechanism to embed in the editor several well-formedness rules required in our DSL specification. For instance, we have added constraints to avoid repeated names, constraints to limit the potential values of some numerical attributes, constraints to restrict

<sup>6</sup><https://github.com/SOM-Research/selfadaptive-IoT-DSL.git>

<sup>7</sup><http://mbeddr.com/>



```

tabular editor for concept Sensor
node cell layout:
partial table {
  horizontal r<> {
    cell Sensor c<"Device"> r<>
    cell { name } c<"ID"> r<>
    cell ( % type % -> { name } ) c<"Type"> r<>
    cell { unit } c<"Unit"> r<>
    cell { threshold } c<"Threshold"> r<>
    cell
      ( / % regions % / ) c<"Regions"> r<>
      /empty cell: <default>
    cell { brand } c<"Brand"> r<>
    cell { communication } c<"Communic."> r<>
    cell ( % gateway % -> { name } ) c<"Gateway"> r<>
    cell ( % topic % -> { name } ) c<"Topic"> r<>
    cell { latitude } c<"Latitude"> r<>
    cell { longitude } c<"Longitude"> r<>
  }
}

```

**Fig. 10:** Definition of the tabular editor for the *Sensor* concept

the potential relationships between nodes, and other constraints that prevent ill-formed models from being built.

## 5.2 Runtime Tool Support

Figure 11 summarizes an operational view of our architecture by distinguishing design time (left-hand side) and runtime (right-hand side).

At design time the user creates an initial IoT system specification model using the DSL. The code generator, presented in Section 5.3, transforms such a specification into a set of deployment and configuration options that describe a MAPE-K loop [9] which is performed at runtime.

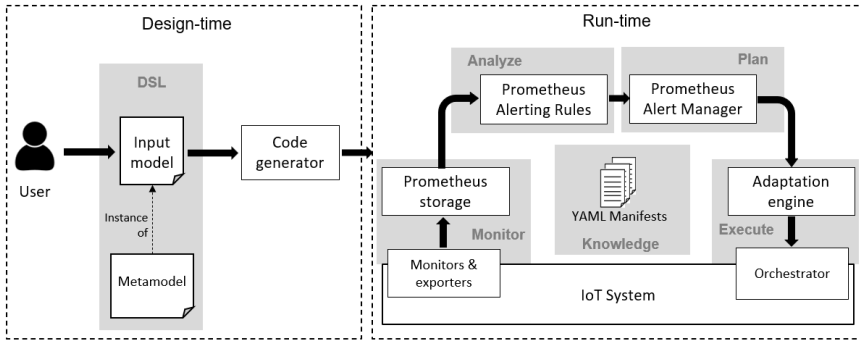
The MAPE-K loop, proposed by IBM for autonomous computing, has been often employed for the design of self-adaptive systems. Indeed, MAPE-K is a reference model to implement adaptation mechanisms in auto-adaptive systems. This model includes four activities (monitor, analyze, plan, and execute) in an iterative feedback cycle that operate on a knowledge base (see Figure 11). These four activities produce and exchange knowledge and information to apply adaptations due to changes in the IoT system.

Based on the MAPE-K loop, our architecture is composed of a set of components and technologies to monitor, analyze, plan, and execute adaptations as illustrated in the right-hand side of Figure 11).

We next describe how our architecture particularizes the generic MAPE-K concepts for self-adaptive IoT systems.

### 5.2.1 Monitor

In the monitoring stage, information about the current state of the IoT system is collected and stored. The collected information is classified into two groups:

**Fig. 11:** Overview

(1) infrastructure and QoS metrics (presented in Table 1); and (2) information that is published in the system's MQTT broker topics such as temperature, humidity, gas levels, and other types of sensor data. These two kinds of information are aligned with the addressed types of events to be detected, i.e., QoS events and sensor events.

We have implemented Prometheus Storage<sup>8</sup> (a time-series database) to store the information collected by the exporters and monitors (such as *kube-state-metrics*<sup>9</sup> and *node-exporter*<sup>10</sup>). Exporters are deployed to convert existing metrics from third-party apps to Prometheus metrics format. The information collected and stored can be queried in real time through the Prometheus user interface or the Grafana dashboard.

**Table 1:** QoS and Infrastructure metrics

Metric	Exporter	Description
Availability	Kube-state-metrics	Equal to 1 if the component being monitored is available, 0 otherwise
CPU	Node Exporter	Number of seconds the CPU has been running in a particular mode
RAM	Node Exporter	Available and total Ram memory of the node
Disk usage	Node Exporter	Available and total disk space of the node
Bandwidth in	Node Exporter	Number of bytes of incoming network traffic to the node
Bandwidth out	Node Exporter	Number of bytes of outgoing network traffic from the node

### 5.2.2 Analyze

The information collected in the monitoring phase must be analyzed, and changes in the system that require adaptations must be identified. To deal

<sup>8</sup><https://prometheus.io/docs/prometheus/latest/storage/>

<sup>9</sup><https://github.com/kubernetes/kube-state-metrics>

<sup>10</sup>[https://github.com/prometheus/node\\_exporter](https://github.com/prometheus/node_exporter)



Remark  
4



Remark  
4 and 23



with this, we have implemented Prometheus Alerting Rules<sup>11</sup> to define alert conditions based on Prometheus query language expressions (PromQL) and to send notifications about firing alerts to the next MAPE-K loop phase. Each IoT system adaptation rule specified through the DSL is transformed into an alert rule of Prometheus.

### 5.2.3 Plan

According to the analysis made in the previous stage, an adaptation plan is generated with the appropriate actions to adapt the system at runtime. The adaptation plan contains the list of actions (scaling, offloading, redeployment, and operate actuator) that the user has defined for each adaptation rule via the DSL. In this stage, Prometheus Alert Manager<sup>12</sup> is used to handle the alerts from the previous stage (Analyze) and routing the adaptation plan to the next stage (Execute). The adaptation plan is sent as an HTTP POST request in JSON format to the configured endpoint (i.e., to the Adaptation Engine).

### 5.2.4 Execute

In the Execute stage, the adaptations are applied to the IoT system following the actions defined in the adaptation plan. To achieve this, we have built the Adaptation Engine, an application developed using Python<sup>13</sup>, flask<sup>14</sup>, and the python API<sup>15</sup> to manage the Kubernetes orchestrator. The adaptation engine can apply two sets of adaptations: (1) architecture adaptations through the orchestrator (e.g., autoscaling an application or offloading a pod); and (2) operating system's actuators to the orchestrator (e.g., by publishing an application or a pod). **Our Adaptation Engine can support the adaptation of any IoT system modeled with our main DSL or even with the extended DSL for the mining industry. There is no need to develop a new Adaptation Engine or modify the existing one for each case study.**

### 5.2.5 Example scenario

As a scenario, we will exemplify the stages of the framework through the adaptation rule specified in Figure 7a. Code for the rule management is automatically generated (Section 5.3), including YAML<sup>16</sup> manifests for deployment, configuration and execution of the monitors, exporters, Prometheus, the Adaptation Engine and other software components implemented in the MAPE-K loop.

In the Monitoring stage, the exporters gather information about CPU consumption of the *fog-f1* node. This information is stored in the Prometheus database. Then, in the Analysis stage, the condition of the adaptation rule is

<sup>11</sup>[https://prometheus.io/docs/prometheus/latest/configuration/alerting\\_rules/](https://prometheus.io/docs/prometheus/latest/configuration/alerting_rules/)

<sup>12</sup><https://prometheus.io/docs/alerting/latest/alertmanager/>

<sup>13</sup><https://www.python.org/>

<sup>14</sup><https://flask.palletsprojects.com/en/2.1.x/>

<sup>15</sup><https://github.com/kubernetes-client/python>

<sup>16</sup>YAML is a data serialization language typically used in the design of configuration files



Remark  
24

verified by executing query expressions in PromQL language. For example, the expression (executed by Prometheus Alerting Rules) that checks if the CPU consumption of the fog-f1 node exceeds 80% for 1 minute is presented in listing 1. Note that we are calculating the average amount of CPU time used excluding the idle time of the node. If the condition is true, the alert signal is sent to the Alert Manager component of the next stage of the cycle (Plan). When the alert is received, the adaptation plan is built containing the two actions (offloading and scaling) and their corresponding information such as container to be offloaded, application to be scaled, number of instances, target nodes and target regions. In the Execute stage, the Adaptation Engine component first performs the Offloading action, and only if it fails, then the second action (Scaling) is performed.

```

1  - alert: ram-consumption
2    expr: 100 - (avg by(node_hostname) (rate(node_cpu_seconds_total{mode=
      "idle",node_hostname="fog-f1"}[15s])) * 100) > 80
3    for: 1m

```

Listing 1: Query expression to check CPU consumption of fog1-f1 node

### 5.3 Code Generator

We have implemented a model-to-text transformation that generates YAML files to deploy the IoT system's container-based applications and the components of each stage of our MAPE-K loop based framework, including its internal logic. For example, we generated the YAML manifest to deploy Prometheus and the rules in PromQL language. More specifically, the generated code includes the following components:

- The container-based IoT applications specified in the input model. Following the running example of Section 2, the YAML manifests for deployment of containers C1, C2, C3, C4, and C5 are generated
- YAML Manifests to deploy the monitoring tools and exporters such as kube-state-metrics, node-exporter, and mqtt-exporter
- YAML code to deploy and configure Prometheus Storage (a time-series database), Prometheus Alerting Rules, and Prometheus Alert Manager. The PromQL code to define the rules (e.g., the code shown in Listing 1) is also generated as a Prometheus configuration file
- YAML manifest to deploy the Adaptation Engine
- and Grafana application to display the monitored data stored in the Prometheus database.

Due to space limitations, we do not show here the model-to-text transformation. You can find all this information in the project repository<sup>17</sup>, including the generated code<sup>18</sup> for our running example from Section 2.

<sup>17</sup><https://github.com/SOM-Research/selfadaptive-IoT-DSL.git>

<sup>18</sup><https://github.com/SOM-Research/selfadaptive-IoT-DSL/tree/main/docs/code-example>



## 6 Empirical Evaluation

Based on the basic methodology for conducting usability studies [18], we have designed and conducted two empirical experiments to validate the expressiveness and ease of use of our DSL: *Experiment 1* focused on mining concepts, and *Experiment 2* focused on architectural concepts. We mostly report on *Experiment 1*, focusing on the DSL mining extension, and give only the highlights of the *Experiment 2* (full details of the latter are available on the project repository<sup>19</sup>). The experiments and their results are outlined below.

### 6.1 Experiment 1: DSL Validation - Mining Concepts

We have designed the first experiment to validate the expressiveness and usability of the DSL regarding the modeling of the mine structure, the control points, sensors and actuators, and adaptation rules to manipulate such actuators (e.g., turn on the mine ventilation system when the methane gas sensor exceeds the threshold value).

#### 6.1.1 Experiment design and setup

Eight subjects participated in the experiment. Participants were experts from the mining domain, but had not been exposed to our DSL before. The goal was to check whether they were able to use it and get their feedback on the experience.

The experiment consisted of an asynchronous screening test (pre-questionnaire) to assess subjects' prior knowledge and suitability for participation, and a synchronous exercise (virtual meeting) with two parts (Sessions 1 and 2). The materials and exercises provided to the participants, the questionnaires and the anonymized answers can be found in the repository of our DSL extended to the mining industry domain<sup>19</sup>.

- Pre-questionnaire (10 min): this screening questionnaire (Q0) was conducted prior to the start of Sessions 1 and 2 to ensure that participants had a basic level of mining knowledge including the structure of underground coal mines, gas monitoring systems, and modeling tools in this domain.
- Session 1 (50 min): In the first 20 minutes of Session 1, we introduced the basic knowledge of IoT systems and the use of the DSL implemented in MPS to model the structure of underground mines, the control points and the IoT devices deployed (sensors and actuators). Next, the participants performed the first modeling exercise about an underground coal mine (with the structure shown in Figure 9), two control points (one at each working face) with three gas sensors and an alarm, a fan, and a control door in the internal tunnel. Each participant was provided with a virtual machine configured with the necessary software to perform the modeling exercise. Finally, the participants filled out a questionnaire (Q1) about the usability and expressiveness of the DSL to model the concepts of the first exercise.

---

<sup>19</sup><https://github.com/SOM-Research/IoT-Mining-DSL>




Remark  
26

- Session 2 (40 min): In Session 2, we first introduced the basic concepts of self-adaptive systems and the design of adaptation rules using our DSL. Next, participants performed the second exercise: modeling three adaptation rules involving sensor data and actuator operation. For example, if any of the methane gas sensors throughout the mine exceed the threshold value for 5 seconds, then turn on the fan and activate the alarms. Finally, participants completed the questionnaire Q2 to report their experience modeling the adaptation rules. Q2 also contained open-ended questions to obtain feedback on the use of the entire tool and suggestions for improvement.

The experiment was conducted in Spanish on three different dates in 2022. The first author of the paper conducted the virtual meetings and ensured that all were equally well executed.

### 6.1.2 Results

Four of the participants were involved in education (either students, teachers, or researchers), while the remaining four were involved in industry. Figure 12 presents the level of general mining knowledge (very low, low, medium, high, and very high) of the eight participants. All of them are aware of the terminology used in the design and structure of underground coal mines. Only two participants were not familiar with cyber-physical or IoT systems for mining. The modeling tools in mining context that they have used are AutoCAD<sup>20</sup> and Minesight<sup>21</sup> for the graphical design of the mine structure, and VentSim<sup>22</sup> for ventilation system simulations. None of them were familiar with MPS.



**Fig. 12:** Participants' mining expertise


Figure 13a presents the responses from questionnaires Q1 and Q2 related to the ease of use of the DSL. Most of them reported that the modeling of the mine structure, the control points, the devices (sensors and actuators),

<sup>20</sup><https://www.autodesk.com/products/autocad>

<sup>21</sup><https://www.ici.edu.pe/brochure/cursos-personalizados/ICI-MINESIGHT-Personalizado.pdf>

<sup>22</sup><https://ventsim.com>

and the adaptation rules were easy. The results are positive and can also be evidenced by the number of right and wrong concepts modeled by the participants (Figure 13b). The number of errors were low (12 of 188 modeled concepts): three incorrect *Rule-conditions* by wrong selection of the unit of measure, four incorrect *Actuators* by wrong assignment of actuator type and location within the mine, three missing *Sensors* not modeled, and two incorrect *Regions* (working faces) whose type was not selected.



**Fig. 13:** Validation results

Through the open-ended questions in questionnaires Q1 and Q2, participants suggested the following improvements to the DSL.

- Include the specification of the coordinates for each region and control points of the mine. Additionally, it would be useful to specify the connection between internal tunnels.
- The condition of an adaptation rule has a single time period. However, it would be useful to associate two time periods for conditions composed of two expressions. For example, the condition  $tempSensorA > 30C(10seconds) \ \&\& \ tempSensorB > 35C(20seconds)$ .
- The mine ventilation system can be activated periodically at the same time each day. It would be useful if the DSL could model adaptation rules whose condition is associated with the time of day.

## 6.2 Experiment 2: DSL Validation - Architectural Concepts

The second experiment aimed to validate the expressiveness and easy of use of concepts more focused on multi-layer architecture modeling: edge, fog, and cloud nodes, container-based applications, and architecture adaptation rules.

The protocol for this experiment was similar to the previous one. Also similarly to the previous experiments, participants reported that the DSL was expressive and easy to use for specifying the multi-layer IoT architecture including



the infrastructure (nodes and servers), container-based applications, and their architectural adaptation rules.

### 6.3 Threats to Validity

Although validation problems in empirical experiments are always possible, we have looked for methods to ensure the quality of the results, analyzing two types of threats: internal and external.

Internal validation concerns factors that could affect the results of the evaluation. To avoid defects in the planning of the experiments and the questionnaires (protocol), all authors of the paper discussed the protocol including the modeling exercises, the dependent variables, and the questions of the questionnaires. In addition, a senior researcher in empirical experiments validated the questionnaires. Another common thread is related to the low number of samples to successfully reveal a pattern in the results. Thirteen users total participated in this empirical validation. Eight participants were involved in Experiment 1, and five different ones in Experiment 2.

External validity addresses threats related to the ability to generalize results to other environments. For example, to validate the population and avoid sampling bias, we conducted a pre-questionnaire to the participants to ensure that they had the necessary basic knowledge and that there was no substantial difference between participants. In addition, at the beginning of each session, we introduced the definition of the concepts required for the experiment. It is important to emphasize that the participants of Experiment 1 were related to the topics of the mining domain, while those of Experiment 2 were computer science researchers.

## 7 Related Work

Modeling of cloud architectures has been widely studied [19–21], including provisioning of resources for cloud applications. Nevertheless, these proposals do not cover multi-layered architectures involving fog and edge nodes. This is also the case for Infrastructure-as-code (IaC) tools.

Regarding IoT system modeling, there is currently no globally accepted metamodel for specifying the architecture and adaptability of IoT systems. Some approaches propose methods to model aspects of the IoT system such as its architecture, functionality, or software deployment [5–7]. Studies such as [5, 6] focus on the modeling of sensors, actuators, and software functionalities, while [7] addresses the modeling of application deployment at the fog layer. Other concepts such as edge nodes, fog nodes, asynchronous communications and self-adaptive capabilities are not supported by these studies. Some other approaches specialize in the modeling of the IoT aspects and propose specific DSLs for that, e.g., [22–24]. However, they do not address the specification and generation of code to deal with runtime adaptations.

Regarding adaptation policies works such as [25–30] propose partial solutions as they either restrict the parts of the system that could be adapted or



Remark  
10 and  
18

offer a limited expressiveness in the definition of the rules (e.g., no trigger condition). In fact, the definition of flexible adaptation rules mixing different types of triggering events remains an open challenge as described in our previous literature review [31]. This was one of the initial motivations of our work.

Works more closely related to our proposal follow. Lee et al. [4] present a self-adaptive framework to dynamically satisfy functional requirements for IoT systems. This framework enables modeling of the IoT environment through a finite-state machine. However, adaptations are only enabled at the device layer level (e.g., opening a window when the temperature exceeds a limit). Rules for other layers or involving several layers are not supported. Garland et al. [32] present Rainbow, a framework for adapting a software system when a constraint is violated. constraints are expressed in terms of performance, cost, or reliability, and the adaptations depend on a set of predefined actions that the control infrastructure can perform. Similarly, Weyns et al. [33] propose MARTA, an architecture-based adaptation approach to automate the management of IoT systems employing runtime models. However, Rainbow and MARTA focus only on architectural adaptations (without addressing container-based applications) and do not cover the specification of multi-layer architectures. Petrovic et al. [34] propose SMADA-Fog, a semantic model-driven approach to deployment and adaptation of container-based applications in Fog computing scenarios. Note that, SMADA-Fog does not allow the specification of complex adaptation rules composed of various conditions and actions. Moreover, grouping nodes and IoT devices according to their location is not possible, forbidding the possibility to apply adaptations on group of nodes belonging to a cluster or a given region.

Concerning the specification of IoT systems in underground mining domain, works such as [35–37] focus on the design or deployment of the device layer (sensors and actuators) of monitoring systems, but do not address the mine structure specification, runtime adaptations, and deployment of containerized applications. There are also graphical modeling tools (such as AutoCAD, VentSim, and MineSight) that specialize in designing the mine layout or simulating its ventilation system, leaving out of their scope the modeling and management of the IoT system.

To sum up, ours is the first proposal that enables the specification, deployment and execution of multi-layer IoT architectures (device, edge, fog, and cloud) and the definition of complex rules covering all layers (and combinations of) involving multiple conditions and actions that can, potentially, engage groups of nodes in the same region or cluster of the IoT system. Moreover, our proposal can be easily extended and specialized in different domains such as underground mining.



Remark  
10 and  
18

## 8 Conclusions and Future Work

We have presented a model-based approach for the specification and runtime execution of multi-layered architectures of IoT systems and their self-adaptation rules. Our approach comprises a new DSL to model such systems, a code generator, and a runtime infrastructure, based on the MAPE-K loop, to monitor and adapt the IoT system at runtime based on a variety of self-adaptive rules, involving architectural adaptations and rules for functional requirements. The full process is assisted by a set of open-source tools that have been released as part of this work. We have also validated the usability and extensibility of the DSL by applying to a particular scenario: coal underground mining.

As part of our future work, we will address the suggestions made by the participants of the performed empirical validations, including a new visual renderer of the modeled architecture to complement the current projections. We also plan to facilitate the definition of complex self-adaptive rules by predefining a set of common patterns such as canary, rolling update, and blue-green deployment strategies that could be directly referenced in the definition of a rule. Moreover, we are also interested to automatically discover potentially useful adaptation rules by analyzing historical log data from the IoT system (e.g., focusing on previous system crashes) with machine learning. Finally, we will create additional extensions to the DSL. In particular, one to model Wastewater Treatment Plants as part of an ongoing project.

## References

- [1] Jiang, Y., Huang, Z., Tsang, D.H.: Challenges and solutions in fog computing orchestration. *IEEE Network* **32**(3), 122–129 (2017)
- [2] Rhayem, A., Mhiri, M.B.A., Gargouri, F.: Semantic web technologies for the internet of things: Systematic literature review. *Internet of Things* **11**, 100206 (2020)
- [3] Brambilla, M., Cabot, J., Wimmer, M.: *Model-driven Software Engineering in Practice*, 2nd Edn. Synthesis Lectures on Software Engineering. Morgan & Claypool Publishers, USA (2017)
- [4] Lee, E., Seo, Y.-D., Kim, Y.-G.: Self-adaptive framework based on mape loop for internet of things. *sensors* **19**(13), 2996 (2019)
- [5] Patel, P., Cassou, D.: Enabling high-level application development for the internet of things. *Journal of Systems and Software* **103**, 62–84 (2015)
- [6] Ciccozzi, F., Spalazzese, R.: Mde4iot: supporting the internet of things with model-driven engineering. In: *Int. Symposium on Intelligent and Distributed Computing*, pp. 67–76 (2016)



- [7] Yigitoglu, E., Mohamed, M., Liu, L., Ludwig, H.: Foggy: a framework for continuous automated IoT application deployment in fog computing. In: IEEE Int. Conf. on AI & Mobile Services, pp. 38–45 (2017)
- [8] Alfonso, I., Garcés, K., Castro, H., Cabot, J.: Modeling self-adaptive IoT architectures. In: 2021 ACM/IEEE Int. Conf. on Model Driven Engineering Languages and Systems Companion (MODELS-C), pp. 761–766 (2021)
- [9] Kephart, J.O., Chess, D.M.: The vision of autonomic computing. Computer **36**(1), 41–50 (2003)
- [10] Latifah, A., Supangkat, S.H., Ramelan, A.: Smart building: A literature review. In: Int. Conf. on ICT for Smart Society (ICISS), pp. 1–6 (2020)
- [11] Al-Qamash, A., Soliman, I., Abulibdeh, R., Saleh, M.: Cloud, fog, and edge computing: A software engineering perspective. In: 2018 International Conference on Computer and Applications (ICCA), pp. 276–284 (2018). IEEE
- [12] Dustdar, S., Avasalcai, C., Murturi, I.: Edge and fog computing: Vision and research challenges. In: 2019 IEEE Int. Conf. on Service-Oriented System Engineering (SOSE), pp. 96–9609 (2019). IEEE
- [13] Yousefpour, A., Fung, C., Nguyen, T., Kadiyala, K., Jalali, F., Niakanlahiji, A., Kong, J., Jue, J.P.: All one needs to know about fog computing and related edge computing paradigms: A complete survey. Journal of Systems Architecture **98**, 289–330 (2019)
- [14] Mansouri, Y., Babar, M.A.: A review of edge computing: Features and resource virtualization. Journal of Parallel and Distributed Computing **150**, 155–183 (2021)
- [15] Gómez, A., Iglesias-Urkia, M., Belategi, L., Mendialdua, X., Cabot, J.: Model-driven development of asynchronous message-driven architectures with asyncapi. Software and Systems Modeling, 1–29 (2021)
- [16] Mishra, B., Kertesz, A.: The use of mqtt in m2m and iot systems: A survey. IEEE Access **8**, 201071–201086 (2020)
- [17] Berger, T., Völter, M., Jensen, H.P., Dangprasert, T., Siegmund, J.: Efficiency of projectional editing: A controlled experiment. In: Proc. of the 24th ACM SIGSOFT Int. Symposium on Foundations of Software Engineering, pp. 763–774 (2016)
- [18] Rubin, J., Chisnell, D.: Handbook of Usability Testing: How to Plan, Design and Conduct Effective Tests. John Wiley & Sons, New Jersey

(2008)

- [19] Sandobalin, J., Insfran, E., Abrahão, S.: ARGON: A model-driven infrastructure provisioning tool. In: ACM/IEEE 22nd Int. Conf. on Model Driven Engineering Languages and Systems Companion (MODELS-C), pp. 738–742 (2019)
- [20] Sledziewski, K., Bordbar, B., Anane, R.: A DSL-based approach to software development and deployment on cloud. In: 24th IEEE Int. Conf. on Advanced Information Networking and Applications, pp. 414–421 (2010)
- [21] Bergmayr, A., Breitenbücher, U., Kopp, O., Wimmer, M., Kappel, G., Leymann, F.: From architecture modeling to application provisioning for the cloud by combining uml and toasca. In: 6th Int. Conf. on Cloud Computing and Services Science, pp. 97–108 (2016)
- [22] Gomes, T., Lopes, P., Alves, J., Mestre, P., Cabral, J., Monteiro, J.L., Tavares, A.: A modeling domain-specific language for IoT-enabled operating systems. In: IECON 2017-43rd Annual Conf. of the IEEE Industrial Electronics Society, pp. 3945–3950 (2017)
- [23] Eterovic, T., Kaljic, E., Donko, D., Salihbegovic, A., Ribic, S.: An internet of things visual domain specific modeling language based on UML. In: 2015 XXV Int. Conf. on Information, Communication and Automation Technologies (ICAT), pp. 1–5 (2015)
- [24] Barriga, J.A., Clemente, P.J., Sosa-Sánchez, E., Prieto, Á.E.: Simulateiot: Domain specific language to design, code generation and execute iot simulation environments. *IEEE Access* **9**, 92531–92552 (2021)
- [25] Cianciaruso, L., di Forenza, F., Di Nitto, E., Miglierina, M., Ferry, N., Solberg, A.: Using models at runtime to support adaptable monitoring of multi-clouds applications. In: 16th Int. Symposium on Symbolic and Numeric Algorithms for Scientific Computing, pp. 401–408 (2014)
- [26] Ferry, N., Chauvel, F., Song, H., Rossini, A., Lushpenko, M., Solberg, A.: Cloudmf: Model-driven management of multi-cloud applications. *ACM Transactions on Internet Technology (TOIT)* **18**(2), 1–24 (2018)
- [27] Chen, W., Liang, C., Wan, Y., Gao, C., Wu, G., Wei, J., Huang, T.: MORE: A model-driven operation service for cloud-based it systems. In: IEEE Int. Conf. on Services Computing, pp. 633–640 (2016)
- [28] Erbel, J., Korte, F., Grabowski, J.: Comparison and runtime adaptation of cloud application topologies based on OCCI. In: The 8th Int. Conf. on Cloud Computing and Services Science, pp. 517–525 (2018)

- [29] Di Menna, F., Muccini, H., Vaidhyanathan, K.: FEAST: a framework for evaluating implementation architectures of self-adaptive IoT systems. In: Proc. of the 37th ACM/SIGAPP Symposium on Applied Computing, pp. 1440–1447 (2022)
- [30] Cámara, J., Muccini, H., Vaidhyanathan, K.: Quantitative verification-aided machine learning: A tandem approach for architecting self-adaptive IoT systems. In: 2020 IEEE Int. Conf. on Software Architecture (ICSA), pp. 11–22 (2020)
- [31] Alfonso, I., Garcés, K., Castro, H., Cabot, J.: Self-adaptive architectures in IoT systems: a systematic literature review. *Journal of Internet Services and Applications* **12**(1), 1–28 (2021)
- [32] Garlan, D., Schmerl, B., Cheng, S.-W.: Software architecture-based self-adaptation. In: *Autonomic Computing and Networking*, pp. 31–55 (2009)
- [33] Weyns, D., Iftikhar, M.U., Hughes, D., Matthys, N.: Applying architecture-based adaptation to automate the management of internet-of-things. In: *European Conf. on Software Architecture*, pp. 49–67 (2018)
- [34] Petrovic, N., Tomic, M.: Smada-fog: Semantic model driven approach to deployment and adaptivity in fog computing. *Simulation Modelling Practice and Theory* **101**, 102033 (2020)
- [35] Porselvi, T., Ganesh, S., Janaki, B., Priyadarshini, K., *et al.*: IoT based coal mine safety and health monitoring system using lorawan. In: 2021 3rd Int. Conf. on Signal Processing and Communication, pp. 49–53 (2021)
- [36] Mishra, P., Kumar, S., Kumar, M., Kumar, J., *et al.*: IoT based multimode sensing platform for underground coal mines. *Wireless Personal Communications* **108**(2), 1227–1242 (2019)
- [37] Alfonso, I., Gómez, C., Garcés, K., Chavarriaga, J.: Lifetime optimization of wireless sensor networks for gas monitoring in underground coal mining. In: 7th Int. Conf. on Computers Communications and Control, pp. 224–230 (2018)

## A Evaluation of System Self-Adaptations

To evaluate the self-adaptation capability of our approach, we conducted three empirical evaluations, one for each type of adaptation (scaling, offloading, and redeployment). For each adaptation assessment we have designed a simple scenario in which an IoT system faces an event that forces adaptations. **The goal of this experiments** is to compare the availability and performance of a non-adaptive IoT system with that of a self-adaptive IoT system that is modeled and managed using our approach. To do so, we have collected and analyzed metrics such as CPU consumption, node availability, and data processing time.

### A.1 Experiment Design and Setup

To test the three architectural adaptations, we have designed a test scenario with the IoT system shown in Figure 10. The system architecture consists of four temperature sensors, two edge nodes (t2.micro AWS instances<sup>23</sup>), one fog node (t2.medium AWS instance), and three applications (*broker-app*, *realtime-app*, and *predictive-app*) executed by four software containers (C1, C2, C3, and C4).

- *broker-app*: MQTT broker that manages messages published by sensor devices. We used Eclipse Mosquitto<sup>24</sup> as message broker because it is lightweight, easy to configure, and is suitable for running at the edge layer. This broker is deployed on the *edge-2* node and executed by the C1 container.
- *realtime-app*: application subscribed to the MQTT broker to consume the data, coming from the sensors, and perform real-time data analysis. For each sensor data received on the broker topics, this application creates a thread or subprocess that performs a series of operations to intentionally generate workload on the node. After processing each sensor data, the result is published in another broker topic.
- *predictive-app*: this application simulates the execution of a predictive algorithm to forecast possible temperature emergencies. The algorithm (subscribed to the broker's topics) receives data from sensors and performs mathematical routines using the NumPy<sup>25</sup> package.

Additionally, we have developed a Python script which publishes random values on topics of the MQTT broker to simulate the four temperature sensors. During the experiment, the sampling rate of the sensors is increased incrementally to perturb the system and induce self-adaptation.

Figure 10 does not show how containers C2, C3, and C4 are deployed on the nodes, because the use of these containers will depend on the type of


---

<sup>23</sup><https://aws.amazon.com/es/ec2/instance-types/t2/>

<sup>24</sup><https://mosquitto.org/>

<sup>25</sup><https://numpy.org/>

adaptation to be tested. Section A.2 shows the protocol of the experiments including the deployment and adaptation of these containers.



**Fig. 10** General test scenario for adaptations

## A.2 Experiment Protocol

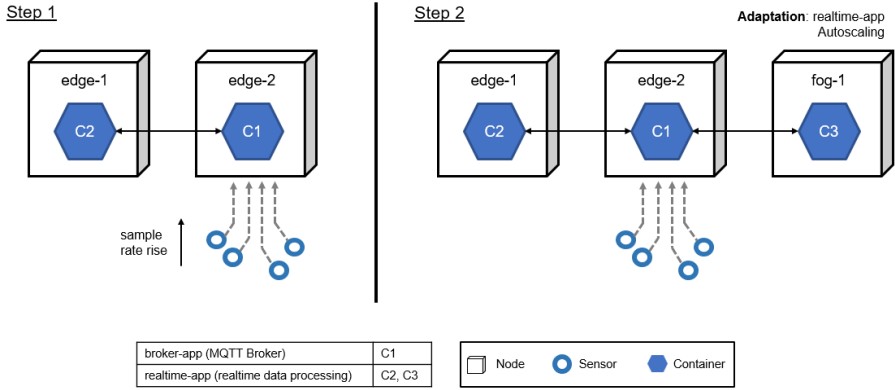
For each type of architectural adaptation we performed an experiment that generally follows the same protocol consisting of the following steps:

1. Model the IoT system (using our DSL) including the adaptation rule to be tested.
2. Run the code generator using the model built in the first step.
3. Deploy the IoT applications and execute our runtime framework using the YAML manifests built by the code generator.
4. Generate disturbances or dynamic events and monitor availability and performance of the system when: (a) the system has no self-adaptation capabilities, and (b) the system self-adapts using our framework.

The variable we manipulate in these experiments (**independent variable**) is the sampling frequency of the sensors in order to generate node overhead, while the variables we monitor (**dependent variables**) are the availability and performance of the nodes. Although the protocol of the three experiments are similar, some aspects change depending on the type of adaptation (scaling, offloading, redeploying) to be tested.

### A.2.1 Scaling an application

The experiment scenario to test *Scaling* adaptation consists of two steps as shown in Figure 11.

**Fig. 11** *Scaling adaptation test scenario*

In step 1, we gradually increase the sampling or monitoring frequency of the sensors to overload the *edge-1* node by increasing the amount of data to be processed by the *C2* container. Initially, we simulate sending 5 data per second to the MQTT broker for two minutes. Then we increase to 12 data per second for the next two minutes. Finally we increase to 30 data per second.

In step 2, after overloading the *edge-1* node, the realtime-app application is scaled. The *Adaptation Engine* of our framework deploys the *C3* container on the *fog-1* node. Then, the data coming from the sensors are distributed for analysis between the *C2* and *C3* containers. This distribution of messages among subscribers is achieved by a balancing strategy known as MQTT shared subscription [16]. In a shared subscription, all clients sharing the same subscription receive messages alternately, i.e., each message will only be sent to one of the subscribed clients.

Figure 12 shows the adaptation rule that we have specified using the DSL to test the *Scalability* adaptation. This rule states that if the CPU usage of the *edge-1* node is greater than 80% for 30 seconds, then scale an instance of the *realtime-app* application on the *fog-1* node. For this adaptation rule, we have chosen to check the CPU usage of the node as it is a metric that reflects the work overhead of the node.

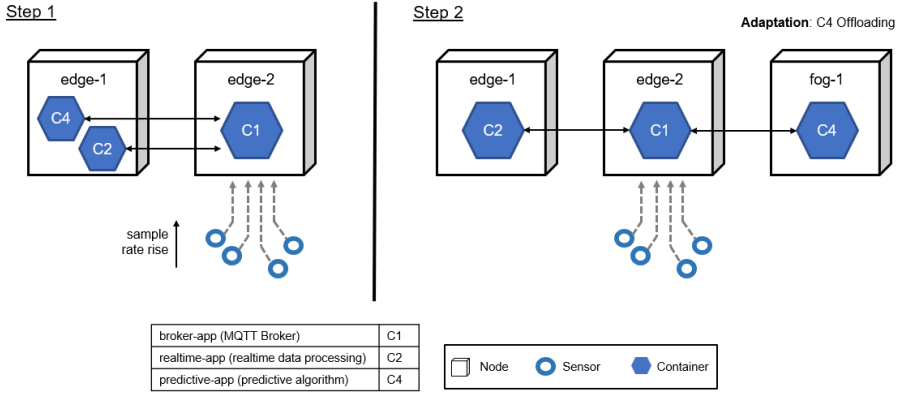
### Scaling App

```
Condition: ( cpu[edge-1] ) > ( 80 % )
Period: 30 s
Actions:
  * Scaling -> Application: realtime-app
                Instances: 1
                Target node(s): fog-1
                Target region(s): << ... >>
                Target cluster: << ... >>
```

**Fig. 12** *Scaling adaptation rule*

### A.2.2 Offloading a container

The experiment scenario to test the *Offloading* adaptation consists of two steps as shown in Figure 13.



**Fig. 13** *Offloading* adaptation test scenario

Similar to our previous experiment, in step 1 the sampling frequency of the sensors is gradually increased (5, 12 and 20 data per second) until an overload is generated in the node *edge-1* due to the increase of sensor data processed by the *C2* container. This node (*edge-1*) hosts the *C2* and *C4* containers that run the *realtime-app* and *predictive-app* applications. Then, in step 2, the system offloads the *C4* container to the *fog-1* node freeing resources on the *edge-1* node.

The adaptation rule modeled is shown in Figure 14: if the CPU usage of node *edge-1* exceeds 80% for 30 seconds, then the *C4* container is offloaded to node *fog-1*.

#### Offloading C4

```

Condition: ( cpu[edge-1] ) > ( 80 % )
Period: 30 s
Actions:
  * Offloading -> Container: C4
    Target node(s): fog-1
    Target region(s): << ... >>
    Target cluster: << ... >>

```

**Fig. 14** *Offloading* adaptation rule

### A.2.3 Redeploying a container

The scenario for testing the *Redeployment* adaptation is the same as shown in Figure 10. First we force a bug in the C1 container by logging into the pod using the command line tool and stopping some system processes. Then, our runtime framework, which constantly monitors the state of containers, detects container unavailability and redeploys it using the *Adaptation Engine*.

The adaptation rule modeled for testing the *Redeployment* of a container is shown in Figure 15. If the container *C1* is detected to be unavailable for more than 20 seconds, then it is redeployed.

```

Redeploying C1
Condition: ( unavailability[C1] )
Period: 20 s
Actions:
    * Redeployment -> Container: C1

```

**Fig. 15** *Redeployment* adaptation rule

## A.3 Results



In the test scenario for each type of adaptation, the system was monitored in two cases: (1) without implementing adaptations, and (2) self-adapting the system according to the configured adaptation rules. The results for each type of adaptation are compared below.

### A.3.1 Scaling an application

Figure 16 presents the CPU usage of the *edge-1* node by increasing the sampling frequency of the sensors. The colored shaded areas represent different sampling frequencies of the sensors. The blue shading indicates that the sensors publish data to the broker at a frequency of 5 data/sec, the green shading 12 data/sec, and the purple shading 30 data/sec. For both cases (non-adaptive and self-adaptive), it was evidenced how the CPU consumption of the node increased when the amount of data to be processed increased too. However, when the CPU usage grew to 100%, the *edge-1* node failed at *tfail* time (Figure 16(a)) for the case of not using adaptations, while implementing the adaptation rule the system auto-scaled the *realtime-app* application (at *tscale* time), the workload was reduced for the *edge-1* node (16(b)), preventing it from failing.

Similarly, Figure 17 shows the time spent by the C2 container to process the data published in the broker by the sensors. For a non-adaptive system (Figure 17(a)) the processing time for some data was reached to grow up to 13.5 sec until the node failed due to work overload. On the other hand, the self-adaptive system (Figure 17(b)) reached processing times of 8.8 sec, then the system auto-scaled the *realtime-app* application at *tscale* time and the data processing time dropped back below 1 sec.



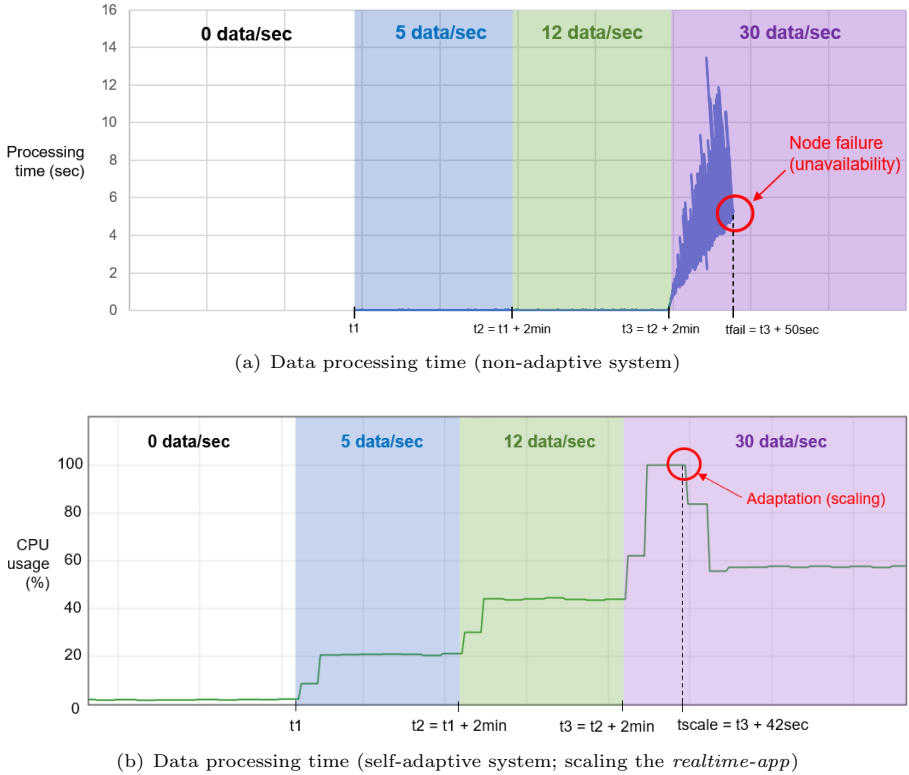


**Fig. 16** *edge-1* node CPU usage; scaling adaptation

### A.3.2 Offloading a container

The CPU consumption of the *edge-1* node is shown in Figure 18. As in the results of the *Scaling* adaptation fitting, the colored shades in the figure represent different data sending frequencies from the sensors. Note that there is a constant CPU usage (43% approx.) before increasing the sampling frequency of the sensors. This CPU usage is caused by *C4* container that simulates the predictive algorithm. As the sampling frequency increased, the CPU consumption of the node also increased. For the non-adaptive system (Figure 18(a)) the node overloaded (CPU usage reached 100%) and failed 26 seconds ( $t_{fail}$ ) after increasing the sampling rate from 12 to 20 data per second. In the case of the self-adaptive system, the CPU consumption of the node did not reach 100% due to the adaptation. The *C4* container was offloaded to the *fog-1* node at  $t_{offload}$  time reducing the workload on the *edge-1* node, preventing it from failing.

Figure 19 shows the processing time of the *C2* container data. This processing time increased considerably when the node used 100% of CPU as is the case of the non-adaptive system (Figure 19(b)), in which the processing time of some data reached 17 seconds before the node failed. On the other hand,





**Fig. 17** Processing time data of *C2* container; scaling adaptation

the self-adaptive system experience processing times of less than 0.3 seconds. This behavior because the *edge-1* node never reached high CPU consumption.


### A.3.3 Redeploying a container

For this Redeployment test, the dependent variable is the availability of the container to be adapted (i.e. to be redeployed). Therefore, we present and analyze the results about the availability of the container.


Figure 20 shows the state (available or unavailable) of the *C4* container for both cases: non-adaptive and self-adaptive systems. In the case of the non-adaptive system, the container is unavailable once its failure has been induced at time  $t1$ . On the contrary case, the self-adaptive system detects that the container is unavailable for 20 seconds and starts the redeployment process. It took approximately 35 seconds to remove the *C4* container and redeploy it. After this procedure, the container changed its status to available again.




**Fig. 18** *edge-1* node CPU usage; offloading adaptation



**Fig. 20** Availability of the *C4* container



(a) Data processing time (non-adaptive system)

(b) Data processing time (self-adaptive system; offloading  $C_4$  container)**Fig. 19** Processing time data of  $C_2$  container; offloading adaptation

## A.4 Threats to Validity

The threats identified in this experiment and how we addressed them are presented below.

**Random irrelevancies in the setting:** some factors outside the experiment (such as network instability or unexpected node failures) could disturb the outcome. In this experiment, we tried to guarantee the stability of the infrastructure using AWS cloud services, which guarantee high availability. Additionally, to avoid disturbances in the delivery of the data sent by sensors, we ran the scripts (simulating the sensors) on an EC2 instance of AWS deployed in the same virtual private cloud as the nodes. This way we guarantee that the data generated in the device layer will be published in the broker on a regular basis and with low latency.

**Measurements of dependent variables should be reliable:** to ensure the reliability of the measurement of dependent variables (e.g., latency and

availability), we have run each experiment at least three times and obtained very similar results. To collect these variables data, we used the monitors and exporters deployed by the framework and consulted the information in the Prometheus time series database.

**Mono-operation bias:** the study should include the analysis of more than one dependent variable. In our experiments, we analyzed various QoS metrics and infrastructure of nodes and containers. For example, we collected and displayed CPU utilization, availability, and data processing latency. Additionally, for the *Redeploying* a container adaptation experiment, in addition to analyzing the availability metric, we also capture the unavailability time of the container while it is redeployed.

**Size of the test scenario:** one of the threats is related to the size of the IoT system to be modeled and tested. Since the objective of this validation was to test the architectural adaptations individually, we proposed a small scenario composed of an IoT system with two edge nodes and one fog node. This scenario was sufficient to model an adaptation rule that allowed testing each adaptation. Nevertheless, we have planned a large scenario (as presented in Section B) to validate the scalability of our approach and the execution of concurrent adaptations.

## B Evaluation of Framework Scalability

To evaluate the ability of our framework to address the growth of concurrent adaptations on an IoT system, we have conducted two experiments: the first one that triggers the firing of simultaneous adaptation rules composed of a single action, and the second one that triggers the firing of simultaneous adaptation rules composed of a group of actions. **The goal of this experiments** is to identify scalability limitations and boundaries to perform concurrent adaptations of our MAPE-K based framework to adapt the IoT system at runtime. The design, protocol, and results of the experiment are presented below.

### B.1 Experiment Design and Setup

To test the scalability of our framework, we have designed a test scenario emulating an IoT environmental monitoring system in three underground coal mines (Figure 21). For simplicity, we assume that the three mines have the same structure (tunnels, work fronts, etc.) and the same monitoring system (nodes and applications). Each mine has ten work fronts<sup>26</sup> constantly monitored to ensure the safety of the workers. The IoT system architecture is composed as follows.

- The **device layer** is composed of several types of sensors deployed at different work fronts. Each work front contains sensors to monitor methane (CH<sub>4</sub>), carbon dioxide (CO<sub>2</sub>), carbon monoxide (CO), Hydrogen sulfide (H<sub>2</sub>S), Sulfur dioxide (SO<sub>2</sub>), nitric oxide (NO), nitrogen dioxide (NO<sub>2</sub>), temperature, and air velocity. The information collected by the sensors is sent to a messaging broker (deployed on *edge-1* for the *Mine 1*).
- The **edge and fog layer** nodes run different applications to detect emergencies, control actuators, store information locally and aggregate information to be sent to the cloud nodes.
- The **cloud layer** nodes run a web application to query historical data about and reports on the environmental status of the three mines. A database is also deployed on one of the cloud nodes to store aggregated data.

To set up the test environment, we provisioned EC2 instances from AWS. Instances t2.micro (1 vCPU and 1 GB of memory) for edge nodes, instances t2.small (1 vCPU and 2 GB of memory) for fog nodes, and instances t2.medium (2 vCPU and 4 GB of memory) for cloud nodes. The collection and sending of data from the sensors was simulated using a script written in Python language.


The containerized applications deployed on the nodes include: an MQTT broker that receives and distributes all sensor data; *stel-app* and *twa-app* check that the values of the monitored gases do not exceed their allowable STEL and TWA value; *temp-app* checks that the temperature at the different work fronts does not exceed the allowable limit value (depending on wind speed); *UI-app* is a user interface for real-time querying of sensor data; *local-database*

---


<sup>26</sup> A large underground coal mine in Colombia could have around ten active work fronts, either mining, advancement, or development.

stores the data locally before being aggregated and sent to the cloud; finally, *web-app* and *cloud-database* enable to store and query the aggregated data in the cloud.

This IoT scenario for underground coal mining is intended to be as close as possible to a real implementation following the rules established in the Colombian mining regulations. For example, the STEL and TWA values are suggested by this decree. However, in this empirical validation, we have developed applications that consume the broker messages and perform some kind of analysis, but do not calculate and check the actual STEL and TWA limit values. This is because the objective of these applications is to generate workload on the nodes to test the scalability limitations of the framework.



Application	Container
broker (MQTT Broker)	C1
*stel-app (realtime gas data analysis)	C2, C5, C8
**twa-app (gas limit value 8h)	C3, C6, C9
temp-app (temperature analysis)	C4, C7, C10
UI-app (user interface)	C11
local-database	C12
web-app	C13
cloud-database	C14



\* STEL value is the permissible limit value for a short exposure time (max. 15 min.)

\*\* TWA value is the permissible limit value for an average time of 8 hours

**Fig. 21** Large mining IoT system to be modeled



### B.1.1 Experiment 1

In Experiment 1, we have tested the activation and execution of simultaneous adaptation rules composed of a single action. We have defined and triggered an adaptation rule for 1, 5, 10, 20, and 30 work fronts (five independent tests).

Figure 22 shows the adaptation rule for *Work Front 1* (a similar rule was specified for each work front): if the CPU consumption of node *edge-3* exceeds 80% for 60 seconds, then offload the container *c3* to node *edge-2*. In this experiment we have chosen the offloading action, which implies more effort for the *Adaptation Engine*, since it involves the removal and creation of a container on a different node.

#### Work Front 1 Rule

```

Condition: ( cpu[edge-3] ) > ( 80 % )
Period: 60 s
Actions:
  * Offloading -> Container: c3
                    Target node(s): edge-2
                    Target region(s): << ... >>

```

**Fig. 22** Adaptation rule for *Work front 1* - single action

### B.1.2 Experiment 2

In Experiment 2, we have defined adaptation rules involving several actions. Similar to Experiment 1, we have defined rules for 1, 5, 10 and 20 work fronts (4 independent tests). Figure 23 shows the adaptation rule for *Work Front 1* (for the other work fronts, the rules are similar), whose list of actions includes one offload, three scaling, and two redeployment. Note that each scaling action is intended to deploy three instances of the application to any node in the mine. This is to detect problems with the execution of concurrent adaptations and limitations in container allocation.

```

Work Front 1
Condition: ( cpu[edge-3] ) > ( 80 % )
Period: 60 s
Actions:
  ☒ all actions
    * Offloading -> Container: c2
                      Target node(s): edge-2
                      Target region(s): << ... >>
    * Scaling -> Application: stel-app
                      Instances: 3
                      Target node(s): << ... >>
                      Target region(s): mine1
    * Scaling -> Application: twa-app
                      Instances: 3
                      Target node(s): << ... >>
                      Target region(s): mine1
    * Scaling -> Application: temp-app
                      Instances: 2
                      Target node(s): << ... >>
                      Target region(s): mine1
    * Redeployment -> Container: c11
    * Redeployment -> Container: c4

```

**Fig. 23** Adaptation rule for *Work front 1* - multiple actions

## B.2 Experiment Protocol

Both Experiment 1 and Experiment 2 follow the same protocol, consisting of the following steps.

1. Model the IoT system (using our DSL) including the adaptation rules to be tested.
2. Run the code generator using the model built in the first step.
3. Deploy the IoT applications and execute our runtime framework using the YAML manifests built by the code generator.
4. Execute the Python script that simulates the generation of sensor data and publishes the messages to the broker. In this way, the necessary workload is generated on the nodes for the adaptation rules to be fired.

The independent variable in both experiments is the frequency of data generation. We set the necessary data frequency (90 samples per minute for each sensor) to trigger the adaptation rules. As for the dependent variables, in Experiment 1 we focused on monitoring the time spent by the system to detect the event, to generate the adaptation plan, and to adapt the system. In Experiment 2, in addition to monitoring the time spent on adaptation, we also focused on analyzing the number of errors performing adaptations and the reasons for them.

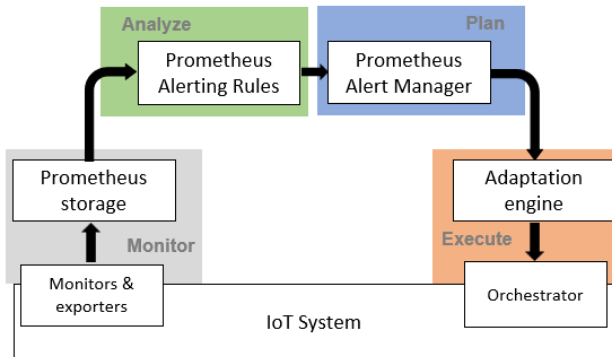
## B.3 Results and Analysis

### B.3.1 Experiment 1

Figure 24 shows the results of the tests performed in this experiment. The information in the table includes the number of adaptation rules configured in the test; the number of errors or failed adaptations; the detection delay refers to the time it takes the system (Prometheus Alerting rules) from the detection of the first rule, to the last one (e.g., for test # 2, the system takes 10,30 seconds from the detection of the event of rule 1 to the detection of the event of rule 5); the time it takes Prometheus Alert Manager to generate the adaptation plan and send it to the Adaptation Engine; and finally, the time it takes the Adaptation Engine to perform the Offloading adaptation, which involves the creation of a new container and the deletion of the old one.

The time spent in the monitoring stage to collect QoS and infrastructure metrics has not been monitored because this is a configurable fixed value for Prometheus. For these experiments, we have set the Prometheus monitoring frequency equal to 4 times every minute (i.e., monitoring every 15 seconds), and the rule evaluation frequency equal to 6 times per minute (i.e., evaluation every 10 seconds). These values were adequate to not generate significant workload on the edge nodes (EC2 t2.micro instances with limited resources).

Test #	Number of rules	Errors	Event detection delay (s)	Adaptation plan design time (s)	Average adaptation time (s)		
					Container creation	Container deletion	Total time
1	1	0	--	1.10	2.06	31.34	33.40
2	5	0	10.30	1.07	2.09	31.32	33.41
3	10	0	12.59	1.12	1.98	31.46	33.44
4	20	0	14.26	1.17	2.19	31.44	33.62
5	30	0	9.96	1.15	2.35	31.40	33.80



**Fig. 24** Experiment 1 results

Findings from the results of this experiment are presented below.

- Ideally the event detection delay (column 4 in Table of Figure 24) should be equal to zero, i.e., all events should be detected at the same time since the increase in CPU consumption was caused at the same time in all nodes. However, there are delays between 10 and 15 seconds approximately, due to the parameters configured in Prometheus (monitoring frequency and rule evaluation frequency). If the monitoring and evaluation frequencies are increased, the event detection delays could be reduced. However, increasing these frequencies would produce significant workload on resource-poor nodes.
- In all cases (even configuring 30 adaptation rules), all actions (offloading) were performed successfully. Approximately one second is required for Prometheus Alert Manager to process an alert, generate the adaptation plan, and send it to the Adaptation Engine. The adaptation time depends on the type of action: the Adaptation Engine, via the K3S orchestrator, takes approximately 2 seconds to create a pod (which hosts a container) and 31 seconds to delete a pod. In this experiment, the MAPE-K components did not fail. However, in the next experiment (Experiment 2) we subjected the framework to more exhaustive tests to detect limitations.
- The average time taken by the adaptation engine to perform a container offload is about 33 seconds, with the removal of the container being the most time consuming task (about 31 seconds). This time is due to the grace period (default 30 seconds) that K3S uses to perform the termination of a pod. When K3S receives the command or API call to terminate a pod, it immediately changes its status to "Terminating" and stops sending traffic to the pod. When the grace period expires, all processes within the pod are killed and the pod is removed. Although our DSL does not currently support grace period configuration, we plan to include the specification of this parameter to ensure safe termination of containers in adaptations that require it (such as offloading or redeployment).

### B.3.2 Experiment 2

In Experiment 2, we set up adaptation rules composed of several actions (as the rule shown in Figure 23). The results obtained are presented in Figure 25, including the number of test rules and actions, the number of failed or unsuccessful actions, the number of nodes that failed, and the average time taken by the Adaptation Engine to perform the successful actions.

Findings from the results of this experiment are presented below.

- For tests 1 and 2 all actions were performed successfully. However, tests 3 and 4 presented failed actions, mainly Scaling type actions (A2, A3, and A4). These scaling actions were not completed, because there were no more resources available on any of the mine nodes to deploy the new container instances. Even some edge nodes (5 for test 3 and 21 for test 4) failed due to work overload causing that some of the A6 actions were also not completed successfully. These results demonstrate that the successful implementation

Test #	Number of rules	Failed actions	Failed nodes	Average time of completed actions (s)					
				A1	A2	A3	A4	A5	A6
1	1 rule 6 actions	0	0	33.46	6.3	6.34	2.17	33.39	32.21
2	5 rules 30 actions	0	0	34.38	6.86	6.32	4.24	32.38	33.22
3	10 rules 60 actions	2 failed A2 actions 7 failed A3 actions 10 failed A4 actions 2 failed A6 actions	5	34.22	6.55	6.45	4.11	32.31	32.34
4	20 rules 120 actions	6 failed A2 actions 14 failed A3 actions 19 failed A4 actions 7 failed A6 actions	21	34.32	6.86	6.92	4.31	32.35	33.31

Actions		
A1. Offload container	A2. Scaling (3 instances of stel-app)	A3. Scaling (3 instances of twa-app)
A4. Scaling (2 instances of temp-app)	A5. Redeploy container	A6. Redeploy container

**Fig. 25** Experiment 2 results

of the adaptations strongly relies on the availability of resources of the target nodes of the actions. In this sense, one of the improvements for our DSL could be to generate warnings to the user when insufficient resources are detected to perform the modeled adaptation rules. In this way we could prevent the implementation of infeasible rules that could fail due to lack of resources.

- The nodes that failed during tests 3 and 4 showed high CPU consumption due to the number of containers assigned to them. Whenever a new pod is deployed on a cluster of nodes (e.g., on one of the edge nodes in mine1), the Kubernetes Scheduler<sup>27</sup> becomes responsible for finding the best node for that pod to run on. Although the Scheduler checks the node resources, it does not analyze the real-time consumption of CPU and Ram memory. Therefore, for Scaling actions in tests 3 and 4, the Scheduler assigned containers to nodes (without checking their current state) causing them to fail. The design and implementation of a Scheduler that analyzes real-time metrics (such as CPU consumption) could avoid this kind of errors.
- The types of actions that require more time to be performed are Offloading and Redeployment. This is because these actions involve the removal of a pod, a task that takes about 31 seconds due to the default grace period set by the K3S orchestrator. On the other hand, the average time it takes to perform the Scalin action depends on the number of instances to be deployed. The Adaptation Engine takes about six seconds to scale three pods or containers, while it takes about four seconds to scale two pods or containers.

<sup>27</sup><https://kubernetes.io/docs/concepts/scheduling-eviction/kube-scheduler/>

## B.4 Threats to Validity

The threat *random irrelevances in the setting* was addressed in the same way as we discussed in Section ?? . Other threats identified in this experiment and how we addressed them are presented below.

**Measurements of dependent variables should be reliable:** to ensure the reliability of the measurement of the dependent variables (e.g., time consumed adapting the system), we performed each experiment at least three times and obtained very similar results. To obtain these metrics, we generated and reviewed log files that record the time and result of each of the tasks performed by the framework components. For example, a log file records the time at which the Adaptation Engine receives the adaptation plan, the time at which each action finishes (including the K3S API command responses), and the result.

**Mono-operation bias:** To avoid this threat, we have planned two experiments involving several tests. Each experiment contains different adaptation rules to analyze constraints on the performance of concurrent adaptations. We increased the number of configured rules and collected more than one dependent variable to analyze the system behavior.

**Size of the test scenario:** The number of nodes (edge, fog, and cloud), sensors, and the size of the mine structure was one of the threats we had to validate. For this, we proposed a scenario with underground coal mines containing 10 work fronts, something usual in medium and large scale mining in the department of Boyacá, Colombia (region that supports part of the doctoral thesis). The simulated sensors covered the generation of data for the variables required by Colombian mining regulations (seven types of gases, temperature, and wind speed). The applications performed the operations that are also required by the regulation. Finally, we provisioned 68 AWS EC2 instances (63 edge nodes, 3 fog nodes, and 2 cloud nodes).